

COM1006 Devices and Networks (Autumn) COM1090 Computer Architectures

Lecture #13

Computer memory

Dr Dirk Sudholt
Department of Computer Science
University of Sheffield

d.sudholt@sheffield.ac.uk

http://staffwww.dcs.shef.ac.uk/~dirk/campus_only/com1006/

Based on Chapters 8, 12, 13 in Clements, Principles of Computer Hardware

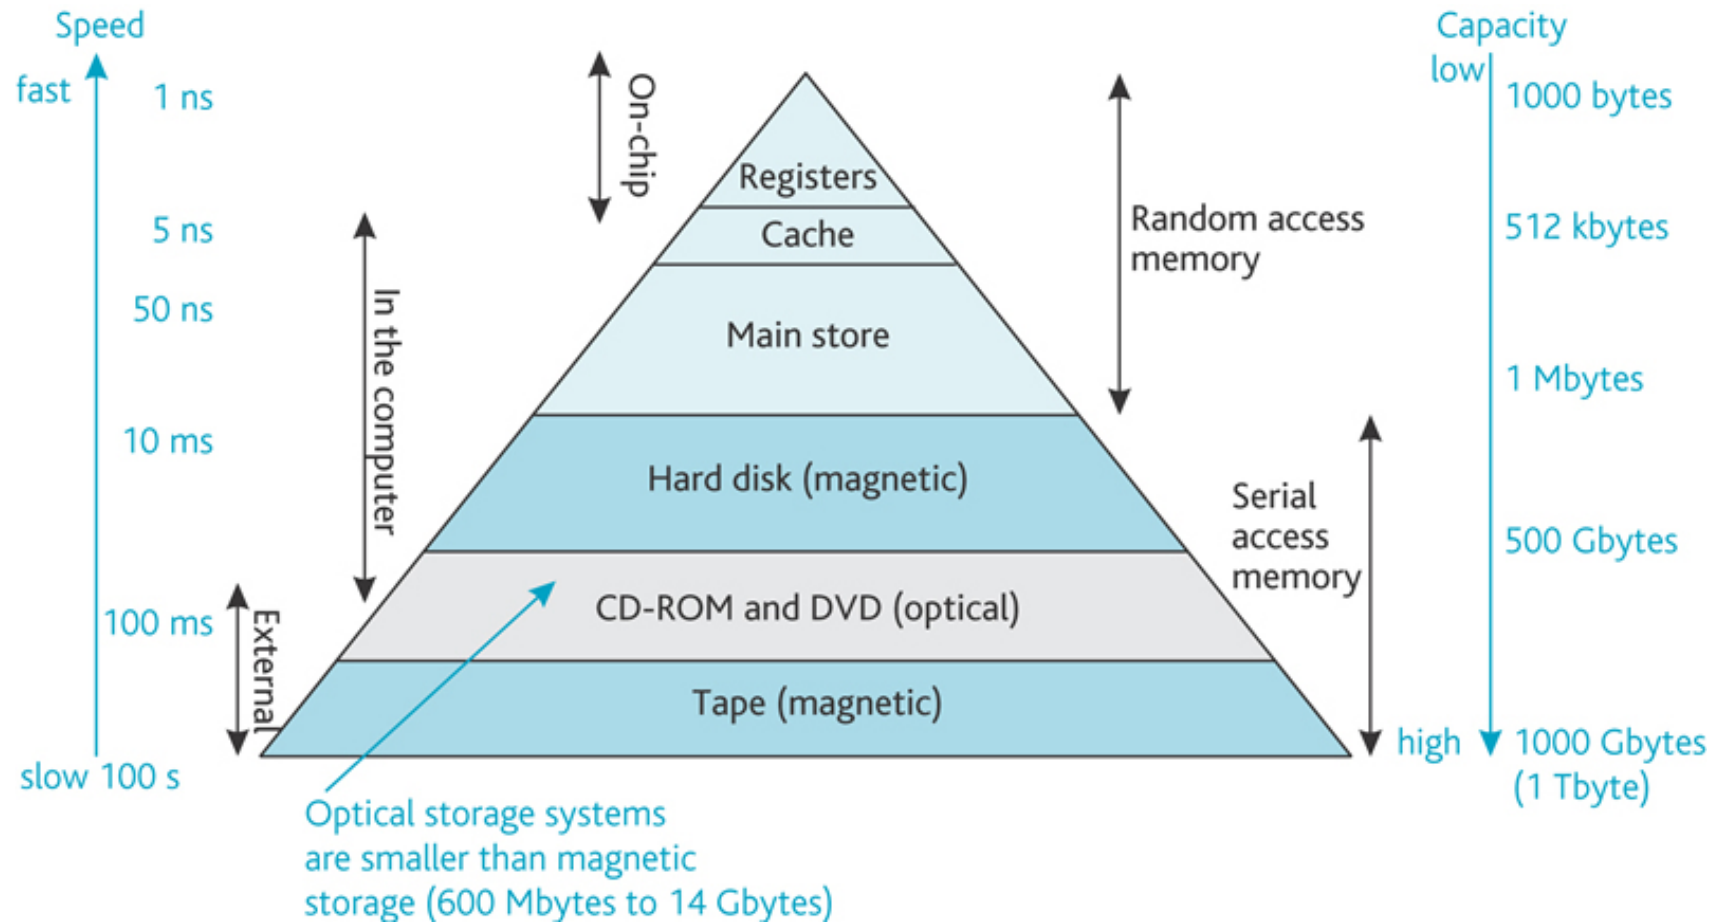
u Aims of this lecture

- To explain the concept of memory hierarchy.
- To explain the way in which cache memory accelerates computer performance.
- To describe the three main strategies for cache memory.
- To explain logical-to-physical mapping in memory management.
- To introduce the concept of virtual memory.
- To point out the similarity between strategies for cache memory and virtual memory.

u Computer memory

- Computer memory systems are not homogeneous.
- Many kinds of memory technology (semiconductor, optical, magnetic) – **why** do we need them?
- ❓ What are the characteristics of ideal computer memory?
 - In practice, memory technologies can be arranged into a **hierarchy** with fast, expensive, low capacity at the top and slow, cheap, large capacity at the bottom.
 - Combining different technologies allows memory systems that are low-cost but close in performance to a system built only with high-speed random-access memory.

Memory hierarchy

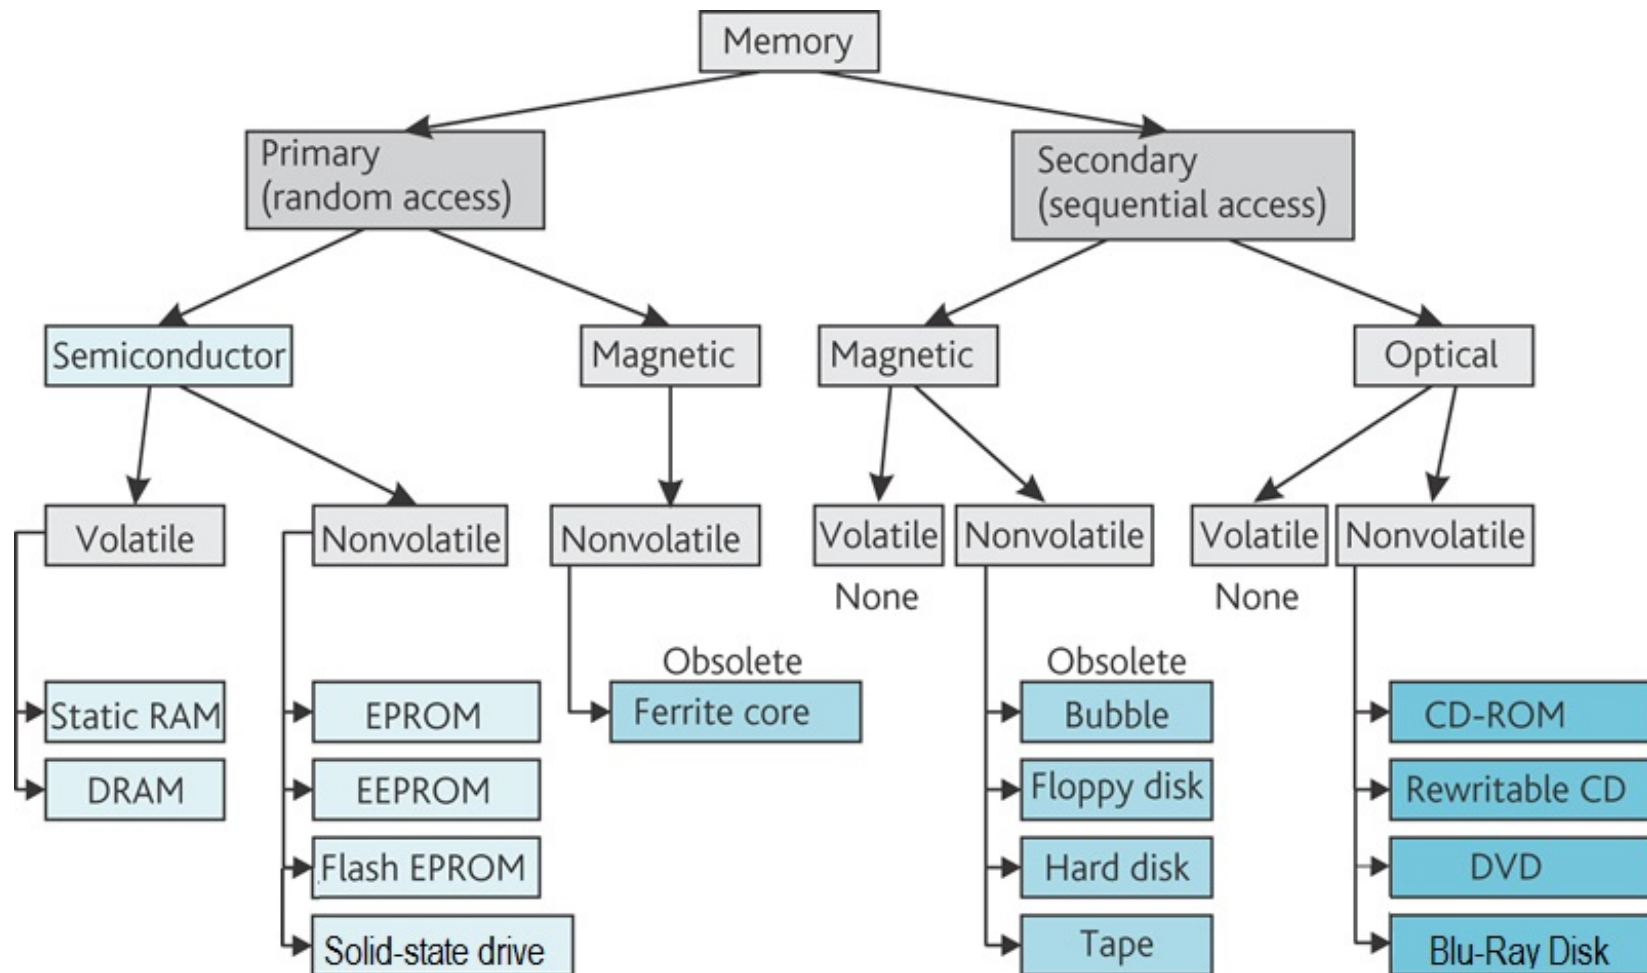


u Memory hierarchy design

- The goal of the hierarchy design is to have the CPU see the entire hierarchy as having the **capacity of the bottom level** while being accessible at the **speed of the top level**.
- Hierarchy is controlled by the hardware (for cache and main store) or by operating system (virtual memory). Use of registers is determined by compiler.
- Hierarchy design is influenced by the phenomenon of **locality of reference**:

The collection of data locations referenced in a short period of time in a running computer often consists of relatively predictable clusters

uClasses of memory



uClasses of memory

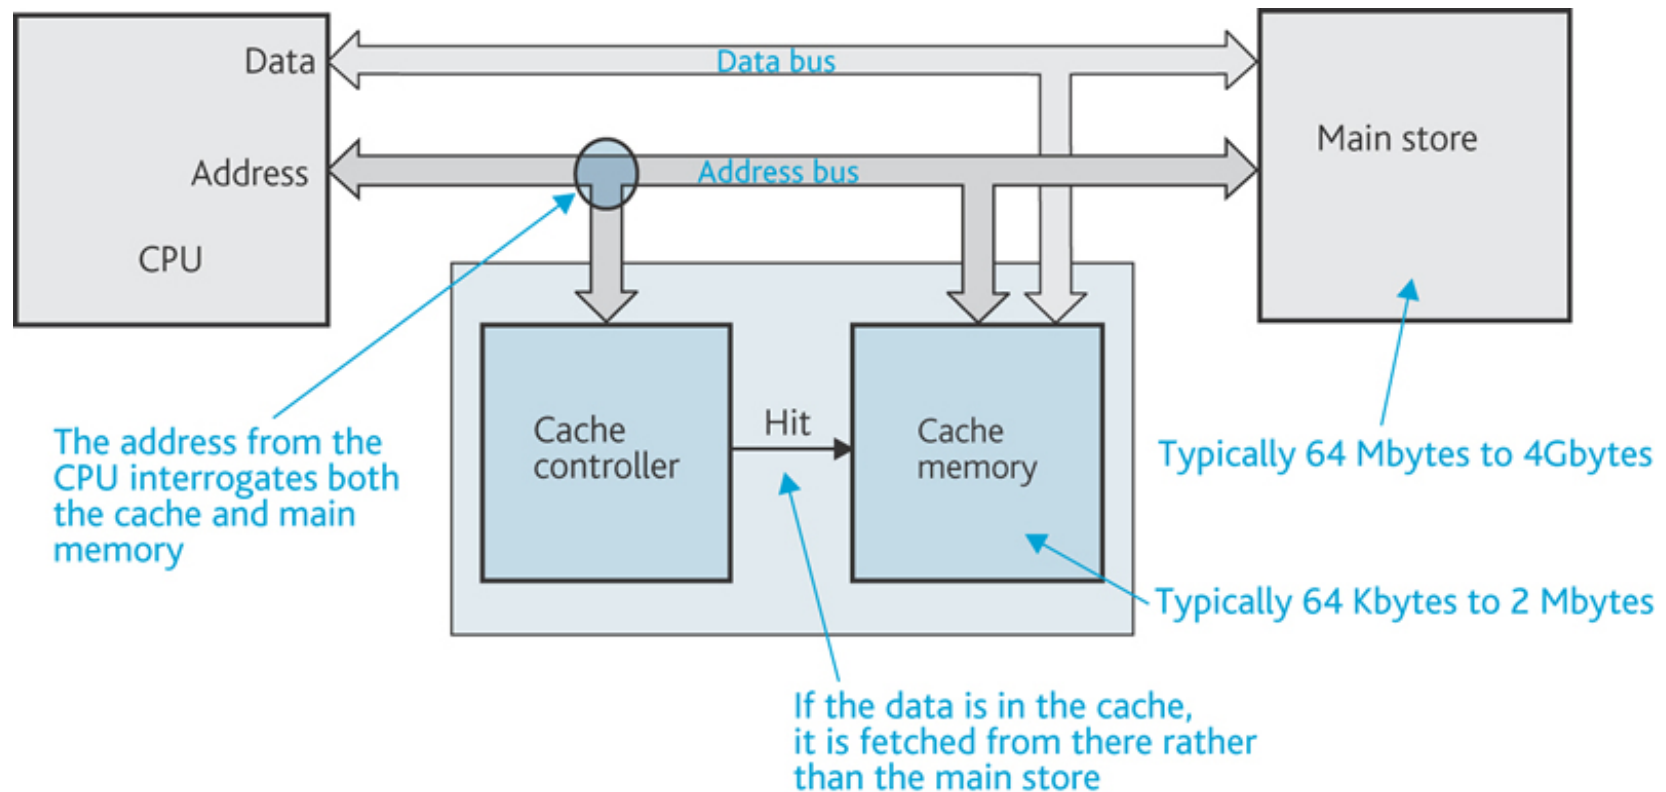
- **Random access memory (RAM)**
 - static (SRAM): based on flip-flops, require power
 - dynamic (DRAM): capacitor stores charge to represent 0/1 state, must be refreshed regularly, compact, low power
- **Read-only memory (ROM)**
 - EPROM: erasable programmable memory, erased by UV light
 - EEPROM: electrically erasable programmable memory
 - Flash EEPROM (and also flash RAM) : nonvolatile memory in which a section of memory cells is erased in a single step (or 'flash'). Basis of many 'solid state' drives



u Cache memory

- Cache memory (from French 'cacher' = 'to hide') can dramatically increase performance at relatively little cost.
- High speed memory, can be accessed rapidly by CPU.
- Analogy: cache is like a personal address book as opposed to telephone directory.
- Cache memory connected in parallel with main memory on address bus (data held in cache is also in main store).
- **Cache controller** determines whether data accessed by CPU is in the cache, or must be obtained from main store.
- Controller returns a **hit** if required data is in the cache.

u Structure of cache memory



u Effect of cache memory on performance

- Key parameter of cache memory is the **hit ratio** (h), defined as ratio of hits to all accesses.
- Hit ratio determined by statistical observation of real system, depends on nature of program being executed and cannot readily be calculated.
- To quantify effect of cache on performance, define the following terms:

Access time of main store	t_m
Access time of cache memory	t_c
Hit ratio	h
Speedup ratio	S

u Cache and performance

- N accesses to system without cache memory requires Nt_m seconds.
- N accesses to system with cache requires $N(ht_c + (1-h)t_m)$ seconds (where $(1-h)$ is the **miss ratio**).
- The figure of merit for a computer with cache is the **speedup ratio**, S , which is ratio of memory access time without cache to access time with cache:

$$S = \frac{Nt_m}{N(ht_c + (1-h)t_m)} = \frac{t_m}{ht_c + (1-h)t_m}$$

u Cache and performance (cont'd)


- Convenient to introduce parameter $k = t_c / t_m$.
- We can now write the expression for S as:

$$S = \frac{t_m / t_m}{ht_c / t_m + (1 - h)t_m / t_m} = \frac{1}{hk + (1 - h)}$$

- Typical values are $t_c = 10$ ns and $t_m = 50$ ns, giving $k = 0.2$.
- Note that $S = 1$ when $h = 0$ (all accesses to main memory).
- When $h = 1$, $S = 1/k$ (all accesses to cache).

u Speedup as a function of hit ratio

- Only when hit ratio near to 90% does the effect of cache become really significant.
- Fortunately effect of locality of reference means that hit ratio is usually very high – data required by processor is highly clustered.




u Cache organisation

- Three main ways of organising cache memory:
 - direct-mapped
 - associative-mapped.
 - set associative-mapped
- Each has different performance/cost trade-offs.
- The approaches differ in how they choose the subset of main memory that is represented in cache.
- The basic unit of data transferred between cache and main store is a **line** (also called a 'block'), usually between 4 and 32 bytes.

u Direct-mapped cache

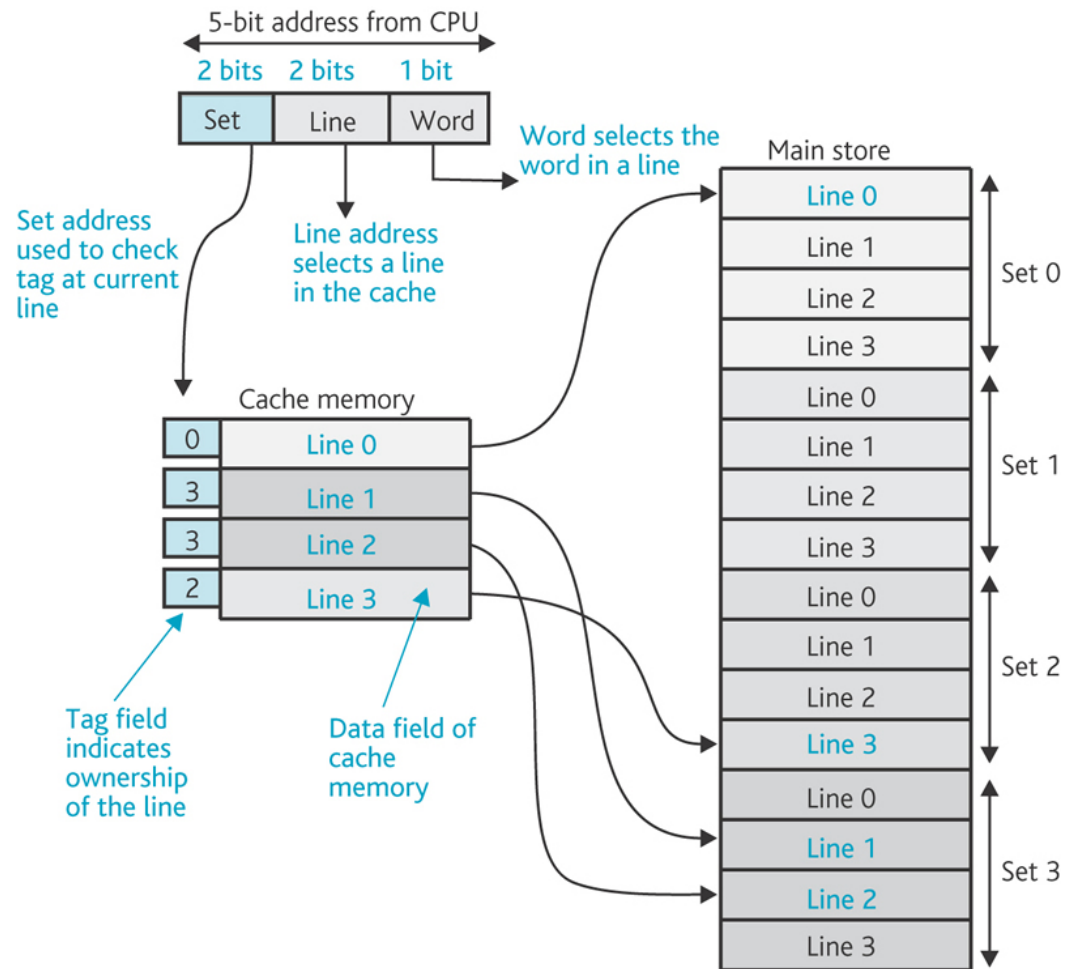
- N Lines are organised into M **sets**. Cache memory contains N/M lines, which may correspond to any set.
- **Example:** Consider main memory of 32 words accessed by a 5-bit address bus from the CPU. Each line contains two words, 4 sets containing 4 lines each ($4 \times 4 \times 2 = 32$)




- Cache is indexed by line, so contains 4 lines in this case.
- **Tag field** in cache records set that the line came from.
- Set in address cross-checked against tag by comparator.

u Direct-mapped cache (cont'd)

- **Example 1:** address 11010_2 is word 0 of line 1 in set 3
- **Hit:** line 1 of cache is used.
- **Example 2:** address 10101_2 is word 1 of line 2 in set 2
- **Miss:** cache is updated by copying contents of address 10101_2 into line 2 of the cache and setting its tag to 2.



u Implementation of direct-mapped cache



u Pros and cons of direct-mapped cache

- Direct-mapped cache is just high-speed memory with built-in comparator, inexpensive and easy to build.
- Cache memory and tag memory are independent and can function in parallel.
- Performance can be poor in some cases. Consider this:

```
repeat
    get data from set  $i$  line  $j$ 
    compare with data from set  $k$  line  $j$ 
until done
```


- Data in line j must be loaded into the cache twice on each iteration of the loop. In practice such cases have little impact on average performance.

u Associative-mapped cache

- Associative-mapped cache places no restrictions on the data it can contain (cf. direct-mapped, which cannot contain the same line from two different sets).
- Address from CPU divided into tag and word fields.
- Example: Consider a system with 1MB (2^{20} bytes) of main store and 64 KB (2^{16} bytes) of associatively-mapped cache.
 - Assume size of line is four 32-bit words (16 bytes)
 - Cache composed of $2^{16}/16 = 4096$ lines
 - Main memory composed of $2^{20}/16 = 65,536 = 2^{16}$ lines
 - Any line in main memory may be stored in cache, so a 16-bit cache tag is required to identify the line.

u Implementing associative-mapped cache

- Associate-mapped cache employs a special **associative memory** that simultaneously compares the address tag to all the tags in cache.
- Hit if input tag matches one of the cache tags, and corresponding line returned from cache.
- Miss if no matching tag in cache.



uPros and cons of associative-mapped cache

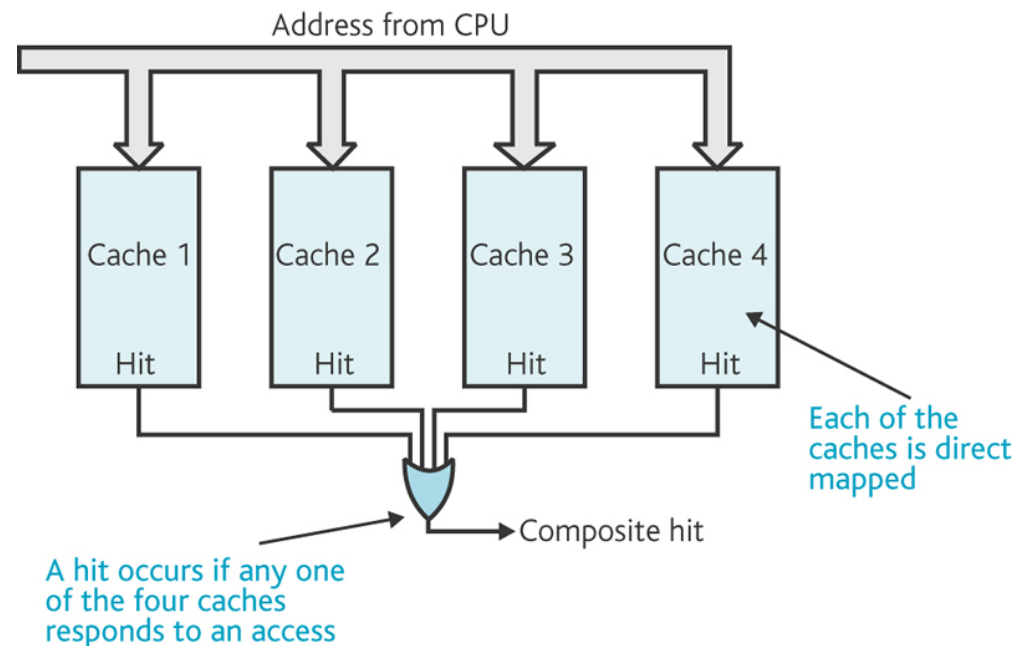
- Associative memory is specialised and is not addressed in the same way as conventional memory:
 - conventional: refer to a specified location
 - associative: do you have this item stored somewhere?
- Large associative memories currently not cost effective.
- When associative cache is full, an existing line must be over-written and this requires a (possibly complex) line **replacement policy**.

u Set-associative mapped cache

- Most computers use a compromise between direct-mapped cache and associative-mapped cache.
- A **set-associative mapped cache** consists of a number of direct-mapped caches operating in parallel.
- If cache has M parallel sets, an M -way comparison is performed on all members of the set.
- Simplest arrangement is two-way set-associative cache in which each line is duplicated.
- This allows line j from set i and line j from set k to be stored at the same time in cache memory (see earlier example).

u Implementation of set-associative cache

- Hit occurs if the requested line and tag occur in any of the caches (4 in this example).
- Logic is not complex because number of caches is small.
- Hit indicated by ORing the hit outputs of all caches.




uMemory management

- In simple computer systems, address generated by CPU corresponds to location in memory.
- Modern operating systems use **memory management** to translate **logical address** generated by CPU to **physical address** (actual memory location).
 - Physical address may correspond to memory contents that are temporarily stored on hard disk, so-called **virtual memory**. Allows running programs **larger than the main memory**.
 - Used in multitasking operating systems to give the impression that each process has sole control of the CPU.
 - Used to protect one process from being corrupted by another process (memory protection).

u Why is logical to physical mapping needed?

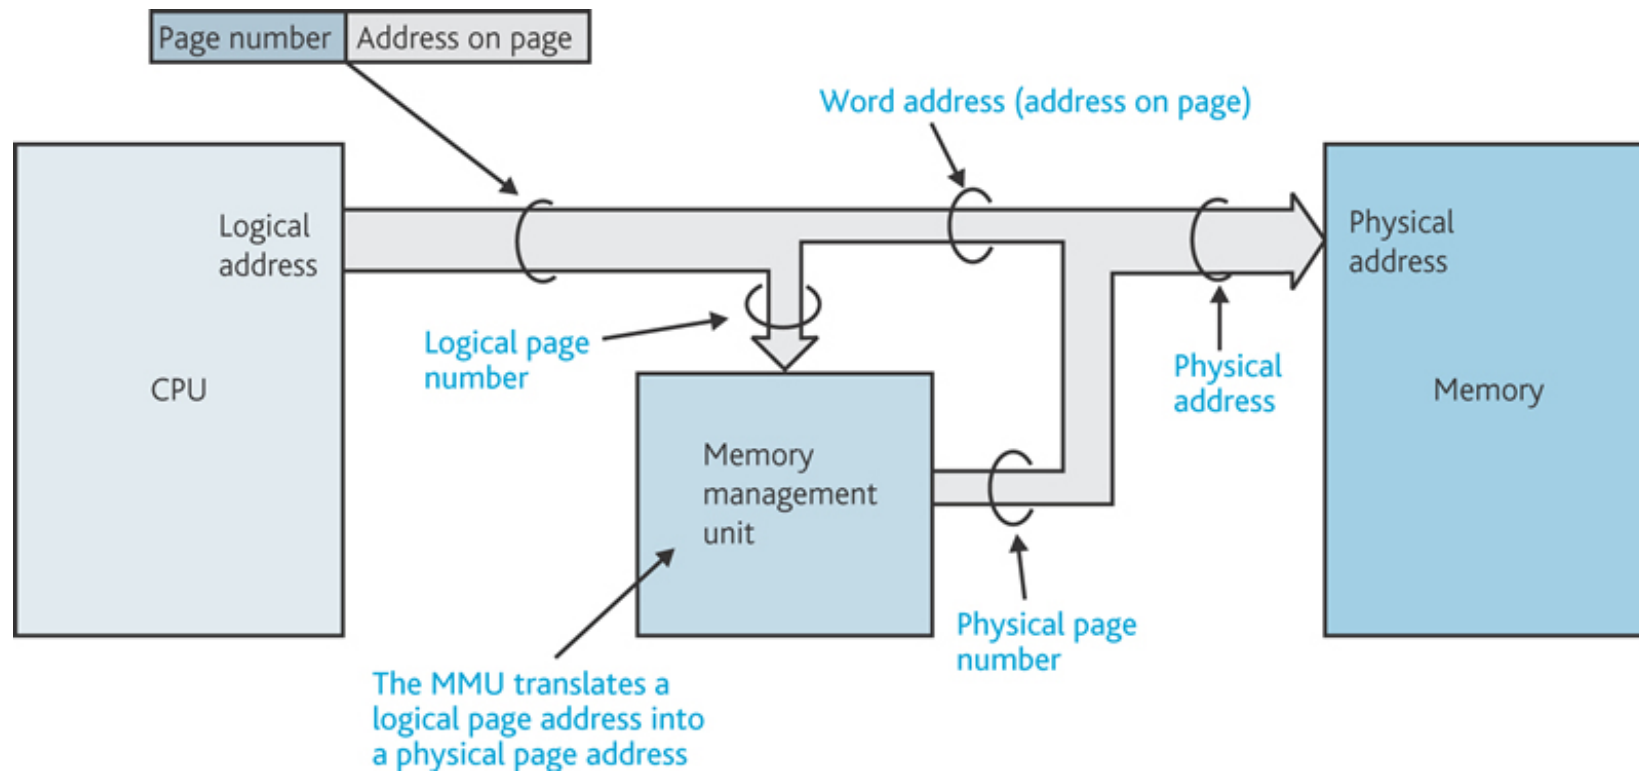
- As the operating system creates new processes and removes old ones, memory becomes **fragmented**.
- Logical address space appears contiguous but physical memory is not.
- If logical address space is larger than physical address space, some memory contents will be temporarily swapped to disk (virtual memory).



u Memory management unit (MMU)

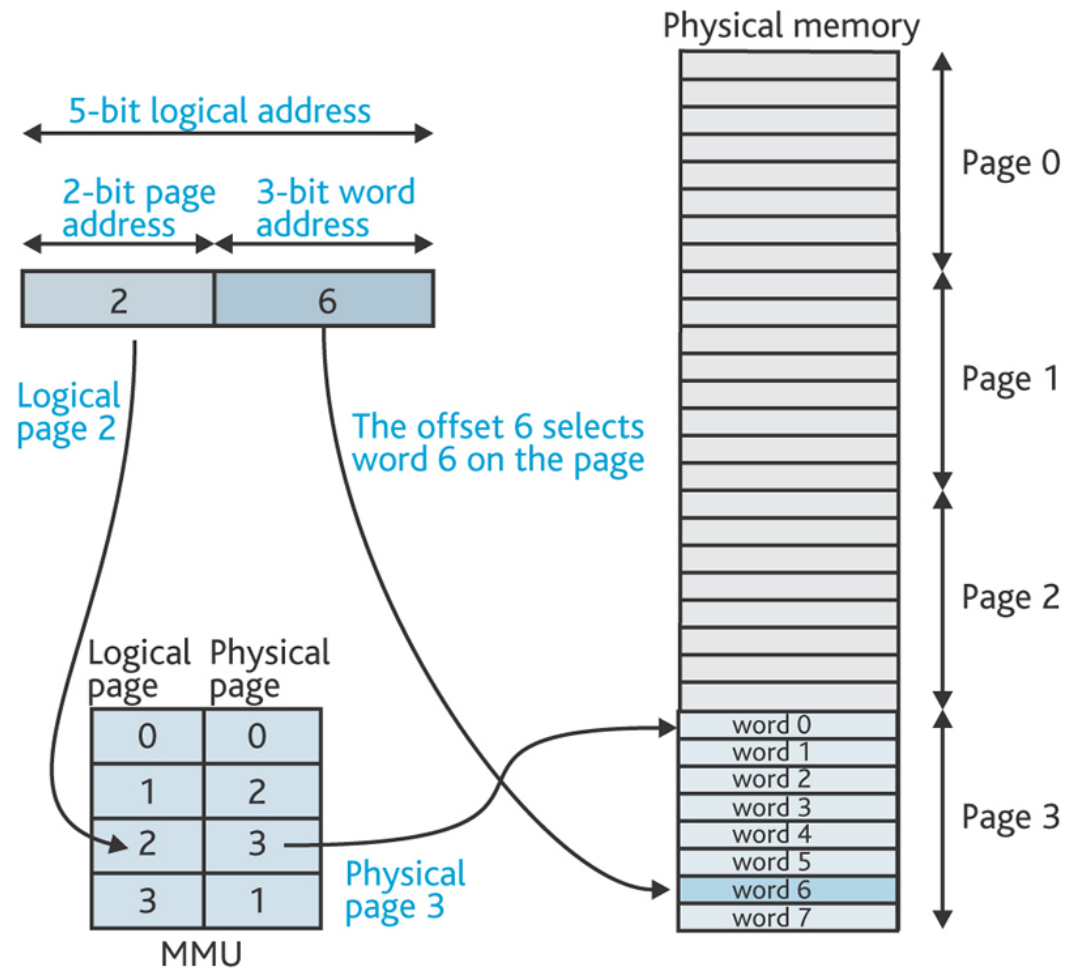
- Modern chips have an integrated **memory management unit** (MMU) (cf. cache controller).
- MMU maps logical address to physical address and keeps track of where data are stored (disk or main store).
- Physical memory is divided into blocks called **pages** (cf. lines).
- Logical address has two parts:
 - page address
 - word address
- Word address goes directly to memory but page address goes to the MMU.

u Memory management unit schematic



u Logical to physical mapping

- Consider 5-bit logical address consisting of 2-bit page address and 3-bit word address.
- MMU maps request for word 6 in logical page 2 to word 6 in physical page 3.



u Dealing with page faults

- Principles governing virtual memory are essentially the same as those governing cache memory.
- If requested page is present in RAM, the logical to physical translation occurs and information is accessed.
- If requested page is not in RAM, then:
 - MMU sends a special interrupt to the processor called a **page fault**.
 - Operating system copies page of data from disc to RAM.
 - Operating system updates page mapping table.
- Only works effectively if data is in RAM most of the time.
- Fortunately the 80:20 rule applies: “80% of the time the processor accesses only 20% of the program”.

uSummary

- The goal of the memory hierarchy is to provide both capacity and speed.
- Three main ways of organising cache memory (direct-mapped, associative-mapped, set associative-mapped) with different performance/cost trade-offs.
- Computers use memory management to translate logical address generated by CPU to physical address.
- In virtual memory, physical memory consists of small high-speed RAM and larger low-speed disc store.
- Similar principles underlie the operation of cache and virtual memory.